

## ex\_nn

December 3, 2021

# 1 Machine Learning: Artificial Neural Networks

Instructions \_\_\_\_\_

This file contains code that helps you get started. You will need to complete the following functions

- predict.m
- sigmoidGradient.m
- randInitializeWeights.m
- nnCostFunction.m

For this exercise, you will not need to change any code in this file, or any other files other than those mentioned above.

## 1.1 Import the required packages

```
[1]: import scipy.io
import numpy as np

from predict import predict
from displayData import displayData
from sigmoidGradient import sigmoidGradient
from randInitializeWeights import randInitializeWeights
from nnCostFunction import nnCostFunction
from checkNNGradients import checkNNGradients
from fmincg import fmincg
```

## 1.2 Setup the parameters you will use for this exercise

```
[2]: input_layer_size = 400;      # 20x20 Input Images of Digits
hidden_layer_size = 25;          # 25 hidden units
num_labels = 10;                 # 10 labels, from 0 to 9
                                   # (note that we have mapped "0" to label 9 to follow
                                   # the same structure used in the MatLab version)
```

## 2 ===== Part 1: Loading and Visualizing Data =====

We start the exercise by first loading and visualizing the dataset. You will be working with a dataset that contains handwritten digits.

### 2.1 Load Training Data


```
[83]: print('Loading and Visualizing Data ...')

mat = scipy.io.loadmat('digitdata.mat')
X = mat['X']
y = mat['y']
y = np.squeeze(y)
m, _ = np.shape(X)

# Randomly select 100 data points to display
sel = np.random.choice(range(X.shape[0]), 49)
sel = X[sel,:]

displayData(sel)
```

Loading and Visualizing Data ...



### 3 ===== Part 2: Loading Parameters =====

In this part of the exercise, we load some pre-initialized neural network parameters.

```
[4]: print('Loading Saved Neural Network Parameters ...')

# Load the weights into variables Theta1 and Theta2
mat = scipy.io.loadmat('debugweights.mat');

# Unroll parameters
Theta1 = mat['Theta1']
Theta1_1d = np.reshape(Theta1, Theta1.size, order='F')
Theta2 = mat['Theta2']
Theta2_1d = np.reshape(Theta2, Theta2.size, order='F')

nn_params = np.hstack((Theta1_1d, Theta2_1d))
```

Loading Saved Neural Network Parameters ...

### 4 ===== Part 3: Implement Predict =====

After training the neural network, we would like to use it to predict the labels. You will now implement the “predict” function to use the neural network to predict the labels of the training set. This lets you compute the training set accuracy.

```
[5]: pred = predict(Theta1, Theta2, X);
print('Training Set Accuracy: ', (pred == y).mean()*100)
```

Training Set Accuracy: 97.52

#### 4.1 Testing (you can skip this block)

To give you an idea of the network’s output, you can also run through the examples one at the a time to see what it is predicting. Run the code in the following block to view examples.

**NOTE:** to avoid the printing of all the sample instances, you can replace *range(m)* with a small number


```
[6]: # Randomly permute examples
rp = np.random.permutation(m)

for i in range(m):
    print(i)
    # Display
    print('Displaying Example Image')
    tmp = np.transpose(np.expand_dims(X[rp[i], :], axis=1))
    displayData(tmp)
```

```
pred = predict(Theta1, Theta2, tmp)
print('Neural Network Prediction: ', pred, '(digit ', pred%10, ')')
```

0


Displaying Example Image



Neural Network Prediction: [6.] (digit [6.] )

1


Displaying Example Image



Neural Network Prediction: [9.] (digit [9.] )

2


Displaying Example Image



Neural Network Prediction: [7.] (digit [7.] )

3


Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

4

Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

5


Displaying Example Image



Neural Network Prediction: [5.] (digit [5.] )

6


Displaying Example Image



Neural Network Prediction: [2.] (digit [2.] )

7

Displaying Example Image



Neural Network Prediction: [8.] (digit [8.] )

8


Displaying Example Image



Neural Network Prediction: [4.] (digit [4.] )

9

Displaying Example Image






Neural Network Prediction: [4.] (digit [4.] )

10

Displaying Example Image



Neural Network Prediction: [6.] (digit [6.] )

11


Displaying Example Image



Neural Network Prediction: [7.] (digit [7.] )

12


Displaying Example Image



Neural Network Prediction: [6.] (digit [6.] )

13


Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

14


Displaying Example Image



Neural Network Prediction: [9.] (digit [9.] )

15

Displaying Example Image



Neural Network Prediction: [4.] (digit [4.] )

16


Displaying Example Image



Neural Network Prediction: [5.] (digit [5.] )

17


Displaying Example Image



Neural Network Prediction: [6.] (digit [6.] )

18


Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

19


Displaying Example Image



Neural Network Prediction: [4.] (digit [4.] )

20


Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

21


Displaying Example Image



Neural Network Prediction: [6.] (digit [6.] )

22


Displaying Example Image



Neural Network Prediction: [7.] (digit [7.] )

23

Displaying Example Image



Neural Network Prediction: [7.] (digit [7.] )

24


Displaying Example Image



Neural Network Prediction: [3.] (digit [3.] )

25

Displaying Example Image





Neural Network Prediction: [1.] (digit [1.] )

26


Displaying Example Image



Neural Network Prediction: [8.] (digit [8.] )

27


Displaying Example Image



Neural Network Prediction: [3.] (digit [3.] )

28


Displaying Example Image



Neural Network Prediction: [9.] (digit [9.] )

29


Displaying Example Image



Neural Network Prediction: [1.] (digit [1.] )

30


Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

31


Displaying Example Image



Neural Network Prediction: [5.] (digit [5.] )

32


Displaying Example Image



Neural Network Prediction: [2.] (digit [2.] )

33


Displaying Example Image



Neural Network Prediction: [8.] (digit [8.] )

34


Displaying Example Image



Neural Network Prediction: [1.] (digit [1.] )

35


Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

36

Displaying Example Image



Neural Network Prediction: [1.] (digit [1.] )

37


Displaying Example Image



Neural Network Prediction: [7.] (digit [7.] )

38


Displaying Example Image



Neural Network Prediction: [10.] (digit [0.] )

39

Displaying Example Image



Neural Network Prediction: [1.] (digit [1.] )

40

Displaying Example Image

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
/var/folders/kj/r4bf7thj1gg5bp153_b6nf740000gn/T/ipykernel_29638/3251442423.py
↳ in <module>
      7     print('Displaying Example Image')
      8     tmp = np.transpose(np.expand_dims(X[rp[i], :], axis=1))
----> 9     displayData(tmp)
      10
      11     pred = predict(Theta1, Theta2, tmp)

~/Desktop/Studium/Period8/MachineLearning/Lab5/Jupyter/Initial code/displayData
↳ py in displayData(data)
      9         ax.imshow(np.transpose(np.reshape(data[i,:], (-1, 20))))
      10         ax.set_axis_off()
---> 11     plt.show()

~/miniforge3/lib/python3.9/site-packages/matplotlib/pyplot.py in show(*args,
↳ **kwargs)
      376     """
      377     _warn_if_gui_out_of_main_thread()
--> 378     return _backend_mod.show(*args, **kwargs)
      379
      380

~/miniforge3/lib/python3.9/site-packages/matplotlib_inline/backend_inline.py in
↳ show(close, block)
      41         display(
      42             figure_manager.canvas.figure,
----> 43             metadata=_fetch_figure_metadata(figure_manager.canvas.
↳ figure)
      44         )
      45     finally:

~/miniforge3/lib/python3.9/site-packages/matplotlib_inline/backend_inline.py in
↳ _fetch_figure_metadata(fig)
      229     if _is_transparent(fig.get_facecolor()):
      230         # the background is transparent
--> 231         ticksLight = _is_light([label.get_color()

      232
      233         for axes in fig.axes
      234         for axis in (axes.xaxis, axes.yaxis)
```



```

~/miniforge3/lib/python3.9/site-packages/matplotlib_inline/backend_inline.py in
↳<listcomp>(.0)
    232                 for axes in fig.axes
    233                 for axis in (axes.xaxis, axes.yaxis)
--> 234                 for label in axis.get_ticklabels()])

    235     if ticksLight.size and (ticksLight == ticksLight[0]).all():
    236         # there are one or more tick labels, all with the same
↳lightness

~/miniforge3/lib/python3.9/site-packages/matplotlib/axis.py in
↳get_ticklabels(self, minor, which)
    1230         if minor:
    1231             return self.get_minorticklabels()
-> 1232         return self.get_majorticklabels()
    1233
    1234     def get_majorticklines(self):

~/miniforge3/lib/python3.9/site-packages/matplotlib/axis.py in
↳get_majorticklabels(self)
    1182     def get_majorticklabels(self):
    1183         """Return this Axis' major tick labels, as a list of `~.text.
↳Text`. """
-> 1184         ticks = self.get_major_ticks()
    1185         labels1 = [tick.label1 for tick in ticks if tick.label1.
↳get_visible()]
    1186         labels2 = [tick.label2 for tick in ticks if tick.label2.
↳get_visible()]

~/miniforge3/lib/python3.9/site-packages/matplotlib/axis.py in
↳get_major_ticks(self, numticks)
    1357         while len(self.majorTicks) < numticks:
    1358             # Update the new tick label properties from the old.
-> 1359             tick = self._get_tick(major=True)
    1360             self.majorTicks.append(tick)
    1361             self._copy_tick_props(self.majorTicks[0], tick)

~/miniforge3/lib/python3.9/site-packages/matplotlib/axis.py in _get_tick(self,
↳major)
    2057         else:
    2058             tick_kw = self._minor_tick_kw
-> 2059         return XTick(self.axes, 0, major=major, **tick_kw)
    2060
    2061     def set_label_position(self, position):

~/miniforge3/lib/python3.9/site-packages/matplotlib/axis.py in __init__(self,
↳*args, **kwargs)

```

```

418
419     def __init__(self, *args, **kwargs):
--> 420         super().__init__(*args, **kwargs)
421         # x in data coords, y in axes coords
422         self.tick1line.set(

~/miniforge3/lib/python3.9/site-packages/matplotlib/_api/deprecation.py in
-> wrapper(*inner_args, **inner_kwargs)
429             else deprecation_addendum,
430             **kwargs)
--> 431         return func(*inner_args, **inner_kwargs)
432
433     return wrapper

~/miniforge3/lib/python3.9/site-packages/matplotlib/axis.py in __init__(self,
-> axes, loc, label, size, width, color, tickdir, pad, labelsizes, labelcolor,
-> zorder, gridOn, tick1On, tick2On, label1On, label2On, major, labelrotation,
-> grid_color, grid_linestyle, grid_linewidth, grid_alpha, **kw)
151         self.apply_tickdir(tickdir)
152
--> 153         self.tick1line = mlines.Line2D(
154             [], [],
155             color=color, linestyle="none", zorder=zorder,
-> visible=tick1On,

~/miniforge3/lib/python3.9/site-packages/matplotlib/lines.py in __init__(self,
-> xdata, ydata, linewidth, linestyle, color, marker, markersize,
-> markeredgewidth, markeredgecolor, markerfacecolor, markerfacecoloralt,
-> fillstyle, antialiased, dash_capstyle, solid_capstyle, dash_joinstyle,
-> solid_joinstyle, pickradius, drawstyle, markevery, **kwargs)
352         self.set_dash_capstyle(dash_capstyle)
353         self.set_dash_joinstyle(dash_joinstyle)
--> 354         self.set_solid_capstyle(solid_capstyle)
355         self.set_solid_joinstyle(solid_joinstyle)
356

KeyboardInterrupt:

```

## 5 ===== Part 4: Sigmoid Gradient =====

Before you start implementing backpropagation, you will first implement the gradient for the sigmoid function. You should complete the code in the sigmoidGradient.m file.

```

[7]: print('Evaluating sigmoid gradient...')
example = np.array([-15, -1, -0.5, 0, 0.5, 1, 15])
g = sigmoidGradient(example)

```

```
print('Sigmoid gradient evaluated at', example, ':')
print(g)
```

Evaluating sigmoid gradient...

```
Sigmoid gradient evaluated at [-15.  -1.  -0.5  0.   0.5  1.  15. ] :
[3.05902133e-07 1.96611933e-01 2.35003712e-01 2.50000000e-01
 2.35003712e-01 1.96611933e-01 3.05902133e-07]
```

## 6 ===== Part 5: Initializing Parameters

To learn a two layer neural network that classifies digits. You will start by implementing a function to initialize the weights of the neural network (randInitializeWeights.py)

```
[8]: print('Initializing Neural Network Parameters ...')

initial_Theta1 = randInitializeWeights(input_layer_size, hidden_layer_size)
initial_Theta2 = randInitializeWeights(hidden_layer_size, num_labels)

# Unroll parameters
initial_Theta1 = np.reshape(initial_Theta1, initial_Theta1.size, order='F')
initial_Theta2 = np.reshape(initial_Theta2, initial_Theta2.size, order='F')
initial_nn_params = np.hstack((initial_Theta1, initial_Theta2))
print(initial_nn_params)
```

Initializing Neural Network Parameters ...

```
[-0.05694027 -0.07043208  0.03860064 ... -0.28053134 -0.21011184
 0.40030573]
```

## 7 ===== Part 6: Implement Backpropagation

Now you will implement the backpropagation algorithm for the neural network. You should add code to nnCostFunction.m to return the partial derivatives of the parameters.

```
[9]: print('Checking Backpropagation...')

# Check gradients by running checkNNGradients
checkNNGradients()
```

Checking Backpropagation...

```
[[-9.27825235e-03]
 [ 8.89911959e-03]
 [-8.36010761e-03]
 [ 7.62813551e-03]
 [-6.74798369e-03]
 [-3.04978931e-06]
```

[ 1.42869450e-05]  
 [-2.59383093e-05]  
 [ 3.69883213e-05]  
 [-4.68759787e-05]  
 [-1.75060084e-04]  
 [ 2.33146356e-04]  
 [-2.87468729e-04]  
 [ 3.35320347e-04]  
 [-3.76215588e-04]  
 [-9.62660640e-05]  
 [ 1.17982668e-04]  
 [-1.37149705e-04]  
 [ 1.53247079e-04]  
 [-1.66560297e-04]  
 [ 3.14544970e-01]  
 [ 1.11056588e-01]  
 [ 9.74006970e-02]  
 [ 1.64090819e-01]  
 [ 5.75736494e-02]  
 [ 5.04575855e-02]  
 [ 1.64567932e-01]  
 [ 5.77867378e-02]  
 [ 5.07530173e-02]  
 [ 1.58339334e-01]  
 [ 5.59235296e-02]  
 [ 4.91620841e-02]  
 [ 1.51127527e-01]  
 [ 5.36967009e-02]  
 [ 4.71456249e-02]  
 [ 1.49568335e-01]  
 [ 5.31542052e-02]  
 [ 4.65597186e-02] ] [[-9.27825236e-03]  
 [ 8.89911960e-03]  
 [-8.36010762e-03]  
 [ 7.62813551e-03]  
 [-6.74798370e-03]  
 [-3.04978914e-06]  
 [ 1.42869443e-05]  
 [-2.59383100e-05]  
 [ 3.69883234e-05]  
 [-4.68759769e-05]  
 [-1.75060082e-04]  
 [ 2.33146357e-04]  
 [-2.87468729e-04]  
 [ 3.35320347e-04]  
 [-3.76215587e-04]  
 [-9.62660620e-05]  
 [ 1.17982666e-04]

```

[-1.37149706e-04]
[ 1.53247082e-04]
[-1.66560294e-04]
[ 3.14544970e-01]
[ 1.11056588e-01]
[ 9.74006970e-02]
[ 1.64090819e-01]
[ 5.75736493e-02]
[ 5.04575855e-02]
[ 1.64567932e-01]
[ 5.77867378e-02]
[ 5.07530173e-02]
[ 1.58339334e-01]
[ 5.59235296e-02]
[ 4.91620841e-02]
[ 1.51127527e-01]
[ 5.36967009e-02]
[ 4.71456249e-02]
[ 1.49568335e-01]
[ 5.31542052e-02]
[ 4.65597186e-02]]

```

The above two columns you get should be very similar.  
 (Left-Numerical Gradient, Right-(Your) Analytical Gradient)

If your backpropagation implementation is correct, then  
 the relative difference will be small (less than 1e-9).

Relative Difference: 2.2957543313343497e-11

## 8 ===== Part 7: Implement Regularization =====

Once your backpropagation implementation is correct, you should now continue to implement the regularization gradient.

```

[10]: print('Checking Backpropagation (w/ Regularization) ... ')

## Check gradients by running checkNNGradients
lambda_value = 3
checkNNGradients(lambda_value)

# Also output the costFunction debugging values
debug_J = nnCostFunction(nn_params, input_layer_size, hidden_layer_size,
                          num_labels, X, y, lambda_value)

print('Cost at (fixed) debugging parameters (w/ lambda = 10): ', debug_J[0][0],

```

```
'(this value should be about 0.576051)')
```

```
Checking Backpropagation (w/ Regularization) ...
```

```
[[-9.27825235e-03]
 [ 8.89911959e-03]
 [-8.36010761e-03]
 [ 7.62813551e-03]
 [-6.74798369e-03]
 [-1.67679797e-02]
 [ 3.94334829e-02]
 [ 5.93355565e-02]
 [ 2.47640974e-02]
 [-3.26881426e-02]
 [-6.01744725e-02]
 [-3.19612287e-02]
 [ 2.49225535e-02]
 [ 5.97717617e-02]
 [ 3.86410548e-02]
 [-1.73704651e-02]
 [-5.75658668e-02]
 [-4.51963845e-02]
 [ 9.14587966e-03]
 [ 5.46101547e-02]
 [ 3.14544970e-01]
 [ 1.11056588e-01]
 [ 9.74006970e-02]
 [ 1.18682669e-01]
 [ 3.81928689e-05]
 [ 3.36926556e-02]
 [ 2.03987128e-01]
 [ 1.17148233e-01]
 [ 7.54801264e-02]
 [ 1.25698067e-01]
 [-4.07588279e-03]
 [ 1.69677090e-02]
 [ 1.76337550e-01]
 [ 1.13133142e-01]
 [ 8.61628953e-02]
 [ 1.32294136e-01]
 [-4.52964427e-03]
 [ 1.50048382e-03]] [[-9.27825236e-03]
 [ 8.89911960e-03]
 [-8.36010762e-03]
 [ 7.62813551e-03]
 [-6.74798370e-03]
 [-1.67679797e-02]
 [ 3.94334829e-02]
 [ 5.93355565e-02]
```

```
[ 2.47640974e-02]
[-3.26881426e-02]
[-6.01744725e-02]
[-3.19612287e-02]
[ 2.49225535e-02]
[ 5.97717617e-02]
[ 3.86410548e-02]
[-1.73704651e-02]
[-5.75658668e-02]
[-4.51963845e-02]
[ 9.14587966e-03]
[ 5.46101547e-02]
[ 3.14544970e-01]
[ 1.11056588e-01]
[ 9.74006970e-02]
[ 1.18682669e-01]
[ 3.81928696e-05]
[ 3.36926556e-02]
[ 2.03987128e-01]
[ 1.17148233e-01]
[ 7.54801264e-02]
[ 1.25698067e-01]
[-4.07588279e-03]
[ 1.69677090e-02]
[ 1.76337550e-01]
[ 1.13133142e-01]
[ 8.61628953e-02]
[ 1.32294136e-01]
[-4.52964427e-03]
[ 1.50048382e-03]]
```

The above two columns you get should be very similar.  
(Left-Numerical Gradient, Right-(Your) Analytical Gradient)

If your backpropagation implementation is correct, then  
the relative difference will be small (less than 1e-9).

Relative Difference: 2.2006042941330916e-11  
Cost at (fixed) debugging parameters (w/ lambda = 10): 0.5760512469501331 (this  
value should be about 0.576051)

## 9 ===== Part 8: Training NN =====

You have now implemented all the code necessary to train a neural network. To train your neural network, we will now use “fmincg”, which is a function which works similarly to “fminunc”. Recall that these advanced optimizers are able to train our cost functions efficiently as long as we provide

them with the gradient computations.

```
[11]: print('Training Neural Network...')

# After you have completed the assignment, change the MaxIter to a larger
# value to see how more training helps.
MaxIter = 150

# You should also try different values of lambda
lambda_value = 1

# Create "short hand" for the cost function to be minimized
y = np.expand_dims(y, axis=1)

costFunction = lambda p : nnCostFunction(p, input_layer_size, hidden_layer_size,
                                         num_labels, X, y, lambda_value)

# Now, costFunction is a function that takes in only one argument (the
# neural network parameters)
[nn_params, cost] = fmincg(costFunction, initial_nn_params, MaxIter)

# Obtain Theta1 and Theta2 back from nn_params
Theta1 = np.reshape(nn_params[0:hidden_layer_size * (input_layer_size + 1)],
                    (hidden_layer_size, (input_layer_size + 1)),
                    order='F')
Theta2 = np.reshape(nn_params[((hidden_layer_size * (input_layer_size + 1))):],
                    (num_labels, (hidden_layer_size + 1)), order='F')
```

Training Neural Network...

Iteration	1	Cost:	[3.49823969]
Iteration	2	Cost:	[3.24347381]
Iteration	3	Cost:	[3.12158557]
Iteration	4	Cost:	[2.97383784]
Iteration	5	Cost:	[2.78142747]
Iteration	6	Cost:	[2.56996023]
Iteration	7	Cost:	[2.49661158]
Iteration	8	Cost:	[1.91191044]
Iteration	9	Cost:	[1.62883088]
Iteration	10	Cost:	[1.41085677]
Iteration	11	Cost:	[1.25879926]
Iteration	12	Cost:	[1.21614708]
Iteration	13	Cost:	[1.09110953]
Iteration	14	Cost:	[1.04490237]
Iteration	15	Cost:	[0.98906894]
Iteration	16	Cost:	[0.92819063]
Iteration	17	Cost:	[0.89113199]
Iteration	18	Cost:	[0.87241911]
Iteration	19	Cost:	[0.8266248]



Iteration	20	Cost:	[0.79457046]
Iteration	21	Cost:	[0.77800163]
Iteration	22	Cost:	[0.72515449]
Iteration	23	Cost:	[0.71050929]
Iteration	24	Cost:	[0.69371792]
Iteration	25	Cost:	[0.67281091]
Iteration	26	Cost:	[0.65998649]
Iteration	27	Cost:	[0.65369317]
Iteration	28	Cost:	[0.63977478]
Iteration	29	Cost:	[0.62722291]
Iteration	30	Cost:	[0.61495667]
Iteration	31	Cost:	[0.59925118]
Iteration	32	Cost:	[0.57943578]
Iteration	33	Cost:	[0.57681014]
Iteration	34	Cost:	[0.57402233]
Iteration	35	Cost:	[0.56448431]
Iteration	36	Cost:	[0.55418161]
Iteration	37	Cost:	[0.54394287]
Iteration	38	Cost:	[0.53291131]
Iteration	39	Cost:	[0.5126614]
Iteration	40	Cost:	[0.49888647]
Iteration	41	Cost:	[0.48317515]
Iteration	42	Cost:	[0.47608402]
Iteration	43	Cost:	[0.47234736]
Iteration	44	Cost:	[0.46786548]
Iteration	45	Cost:	[0.46407074]
Iteration	46	Cost:	[0.46093942]
Iteration	47	Cost:	[0.45626854]
Iteration	48	Cost:	[0.45153547]
Iteration	49	Cost:	[0.44902641]
Iteration	50	Cost:	[0.4476314]
Iteration	51	Cost:	[0.44678635]
Iteration	52	Cost:	[0.44415176]
Iteration	53	Cost:	[0.44240665]
Iteration	54	Cost:	[0.44192092]
Iteration	55	Cost:	[0.44055679]
Iteration	56	Cost:	[0.43923347]
Iteration	57	Cost:	[0.43750361]
Iteration	58	Cost:	[0.43572285]
Iteration	59	Cost:	[0.43027741]
Iteration	60	Cost:	[0.42486924]
Iteration	61	Cost:	[0.41320797]
Iteration	62	Cost:	[0.40430599]
Iteration	63	Cost:	[0.39447605]
Iteration	64	Cost:	[0.39145467]
Iteration	65	Cost:	[0.38963542]
Iteration	66	Cost:	[0.38766581]
Iteration	67	Cost:	[0.38531631]

Iteration	68	Cost:	[0.38291395]
Iteration	69	Cost:	[0.38147827]
Iteration	70	Cost:	[0.38084556]
Iteration	71	Cost:	[0.37911034]
Iteration	72	Cost:	[0.37774182]
Iteration	73	Cost:	[0.37705451]
Iteration	74	Cost:	[0.37618309]
Iteration	75	Cost:	[0.37581191]
Iteration	76	Cost:	[0.37570479]
Iteration	77	Cost:	[0.37526859]
Iteration	78	Cost:	[0.37488594]
Iteration	79	Cost:	[0.3747085]
Iteration	80	Cost:	[0.3745231]
Iteration	81	Cost:	[0.37308652]
Iteration	82	Cost:	[0.37176486]
Iteration	83	Cost:	[0.36994506]
Iteration	84	Cost:	[0.36839749]
Iteration	85	Cost:	[0.36761091]
Iteration	86	Cost:	[0.36593911]
Iteration	87	Cost:	[0.36428952]
Iteration	88	Cost:	[0.36209676]
Iteration	89	Cost:	[0.36131615]
Iteration	90	Cost:	[0.36098567]
Iteration	91	Cost:	[0.36058148]
Iteration	92	Cost:	[0.36014886]
Iteration	93	Cost:	[0.35967547]
Iteration	94	Cost:	[0.35931384]
Iteration	95	Cost:	[0.35882235]
Iteration	96	Cost:	[0.35813908]
Iteration	97	Cost:	[0.35765274]
Iteration	98	Cost:	[0.35687305]
Iteration	99	Cost:	[0.35589518]
Iteration	100	Cost:	[0.35570832]
Iteration	101	Cost:	[0.35560143]
Iteration	102	Cost:	[0.35548232]
Iteration	103	Cost:	[0.35517304]
Iteration	104	Cost:	[0.3543256]
Iteration	105	Cost:	[0.35321992]
Iteration	106	Cost:	[0.3528937]
Iteration	107	Cost:	[0.35257076]
Iteration	108	Cost:	[0.35223261]
Iteration	109	Cost:	[0.35187511]
Iteration	110	Cost:	[0.35168459]
Iteration	111	Cost:	[0.35164941]
Iteration	112	Cost:	[0.35150996]
Iteration	113	Cost:	[0.35136733]
Iteration	114	Cost:	[0.35129319]
Iteration	115	Cost:	[0.35123084]

```

Iteration 116 | Cost: [0.35086671]
Iteration 117 | Cost: [0.35058053]
Iteration 118 | Cost: [0.35048007]
Iteration 119 | Cost: [0.35033374]
Iteration 120 | Cost: [0.35020157]
Iteration 121 | Cost: [0.35013218]
Iteration 122 | Cost: [0.35008993]
Iteration 123 | Cost: [0.34979006]
Iteration 124 | Cost: [0.34905649]
Iteration 125 | Cost: [0.34827103]
Iteration 126 | Cost: [0.3474595]
Iteration 127 | Cost: [0.34576245]
Iteration 128 | Cost: [0.34467584]
Iteration 129 | Cost: [0.34325712]
Iteration 130 | Cost: [0.34214393]
Iteration 131 | Cost: [0.34155858]
Iteration 132 | Cost: [0.34135353]
Iteration 133 | Cost: [0.34085519]
Iteration 134 | Cost: [0.34074223]
Iteration 135 | Cost: [0.34048641]
Iteration 136 | Cost: [0.34021976]
Iteration 137 | Cost: [0.34004822]
Iteration 138 | Cost: [0.3398439]
Iteration 139 | Cost: [0.33921049]
Iteration 140 | Cost: [0.33887219]
Iteration 141 | Cost: [0.33879885]
Iteration 142 | Cost: [0.33847175]
Iteration 143 | Cost: [0.33832326]
Iteration 144 | Cost: [0.33828464]
Iteration 145 | Cost: [0.33823076]
Iteration 146 | Cost: [0.33819064]
Iteration 147 | Cost: [0.33804039]
Iteration 148 | Cost: [0.33794879]
Iteration 149 | Cost: [0.33790889]
Iteration 150 | Cost: [0.33782227]

```

## 10 ===== Part 9: Visualize Weights =====

You can now “visualize” what the neural network is learning by displaying the hidden units to see what features they are capturing in the data.


```

[12]: print('\nVisualizing Neural Network... \n')

displayData(Theta1[:, 1:])

```

## Visualizing Neural Network...



## 11 ===== Part 10: Predicting with learned weights =====

After training the neural network, we would like to use it to predict the labels. The already implemented “predict” function is used by neural network to predict the labels of the training set. This lets you compute the training set accuracy.

```
[13]: pred = predict(Theta1, Theta2, X)
      pred = np.expand_dims(pred,axis=1)
      print('Training Set Accuracy: ', (pred == y).mean()*100)
```

Training Set Accuracy: 99.11999999999999

## 12 Report

**12.1 Note:** This report will not go into too much detail about the implementation, but I did my best to provide comments in the source code to explain what I am doing

### 12.1.1 The Neural Network:

The proposed and initialized neural network consists of 400 (+ 1 bias) input nodes, 25 (+1 bias) hidden nodes in one hidden layer and 10 output nodes (corresponding to the labels from 0 to 9 / 1 to 10). The network is fully connected on all layers and therefore has 10426 weights between

the first and second and 260 weights between the second and third layer. This makes a total of 10686 parameters to adjust. The optimization algorithm in use is vanilla backpropagation, which is optimized using the finncg optimizer. Furthermore, the weights are initialized using Xavier initialization, where we calculate an epsilon value as  $\epsilon = \sqrt{6} / \sqrt{L_{in} + L_{out}}$  and then for each weight randomly sample from a uniform distribution over  $[-\epsilon, +\epsilon]$  and set the corresponding weight to that sampled value. This helps / tries to prevent exploding/vanishing gradients when training.

### 12.1.2 Impact of specific variables:

#### - Lambda

```
[59]: %%capture
from sklearn.model_selection import train_test_split
print('Loading and data ...')
mat = scipy.io.loadmat('digitdata.mat')
X = mat['X']
y = mat['y']

X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.34,
    ↪random_state=10)
print('Training Neural Network...')
#Lambda values:
lambda_values = [1, 7, 70]
#Accuracies
accuracies = {}
training_accuracy = {}
#Cost list
cost_list = []
# After you have completed the assignment, change the MaxIter to a larger
# value to see how more training helps.
MaxIter = 150
#Training
for lambda_value in lambda_values:
    print('Initializing Neural Network Parameters ...')

    initial_Theta1 = randInitializeWeights(input_layer_size, hidden_layer_size)
    initial_Theta2 = randInitializeWeights(hidden_layer_size, num_labels)

    # Unroll parameters
    initial_Theta1 = np.reshape(initial_Theta1, initial_Theta1.size, order='F')
    initial_Theta2 = np.reshape(initial_Theta2, initial_Theta2.size, order='F')
    initial_nn_params = np.hstack((initial_Theta1, initial_Theta2))

    # Create "short hand" for the cost function to be minimized
    costFunction = lambda p : nnCostFunction(p, input_layer_size,
    ↪hidden_layer_size,
                                num_labels, X_train, Y_train,
    ↪lambda_value)
```

```

# Now, costFunction is a function that takes in only one argument (the
# neural network parameters)
[nn_params, cost] = fmincg(costFunction, initial_nn_params, MaxIter)

# Obtain Theta1 and Theta2 back from nn_params
Theta1 = np.reshape(nn_params[0:hidden_layer_size * (input_layer_size + 1)],
                    (hidden_layer_size, (input_layer_size + 1)),
↪order='F')
Theta2 = np.reshape(nn_params[((hidden_layer_size * (input_layer_size +
↪1)))):],
                    (num_labels, (hidden_layer_size + 1)),
↪order='F')

#Predicting
pred = predict(Theta1, Theta2, X_test)
pred = np.expand_dims(pred,axis=1)
accuracy = (pred == Y_test).mean()*100
accuracies[lambda_value] = accuracy
pred = predict(Theta1, Theta2, X_train)
pred = np.expand_dims(pred,axis=1)
training_accuracy[lambda_value] = (pred == Y_train).mean()*100
print('Testing Set Accuracy: ', accuracy)
print('Training Set Accuracy: ', training_accuracy[lambda_value])
print("Final Cost: ", cost[-1])
cost_list.append(cost)


```

```

[61]: import matplotlib.pyplot as plt
xs, ys = zip(*accuracies.items())
plt.plot(xs, ys, "o-")
plt.xlabel("Lambda_value")
plt.ylabel("Accuracy")
xs, ys = zip(*training_accuracy.items())
plt.plot(xs, ys, "ro-")
plt.legend(["testing_accuracy", "training_accuracy"])

```


[61]: <matplotlib.legend.Legend at 0x154165be0>



```
[62]: for cost_arr in cost_list:
        plt.plot(range(len(cost_arr)), cost_arr, "-")

plt.xlabel("iteration")
plt.ylabel("Cost")
plt.legend(["lambda = 1", "lambda = 7", "lambda = 70"])
```

[62]: <matplotlib.legend.Legend at 0x1524ee1f0>



As one can see, the lambda value influences both the accuracy and the minimum cost. For large values of lambda, the test accuracy drops, but the training accuracy does as well. Also, the minimum cost is proportional to the value of lambda, meaning that it is bigger when lambda is bigger. However, this is to be expected, since the use of regularization is to try and prevent overfitting. So what we are trying to achieve is to find a good tradeoff between training and testing accuracy, for example at lambda = 1, the training accuracy is very close to 100, while the test accuracy is a little below that. One might suspect overfitting, even though it is not quite there yet, but with more iterations that effect could be amplified, as seen in the next experiment.

#### - Iteration number

```
[64]: %%capture
from sklearn.model_selection import train_test_split
print('Loading and data ...')
mat = scipy.io.loadmat('digitdata.mat')
X = mat['X']
y = mat['y']

X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.34,
    random_state=10)
print('Training Neural Network...')
#Lambda values:
lambda_values = 7
#Accuracies
accuracies = {}
training_accuracy = {}
```



```

#Cost list
cost_list = []
# After you have completed the assignment, change the MaxIter to a larger
# value to see how more training helps.
iterations = [100, 150, 200]
#Training
for MaxIter in iterations:
    print('Initializing Neural Network Parameters ...')

    initial_Theta1 = randInitializeWeights(input_layer_size, hidden_layer_size)
    initial_Theta2 = randInitializeWeights(hidden_layer_size, num_labels)

    # Unroll parameters
    initial_Theta1 = np.reshape(initial_Theta1, initial_Theta1.size, order='F')
    initial_Theta2 = np.reshape(initial_Theta2, initial_Theta2.size, order='F')
    initial_nn_params = np.hstack((initial_Theta1, initial_Theta2))

    # Create "short hand" for the cost function to be minimized
    costFunction = lambda p : nnCostFunction(p, input_layer_size,
↪hidden_layer_size,
                                num_labels, X_train, Y_train,
↪lambda_value)

    # Now, costFunction is a function that takes in only one argument (the
    # neural network parameters)
    [nn_params, cost] = fmincg(costFunction, initial_nn_params, MaxIter)


    # Obtain Theta1 and Theta2 back from nn_params
    Theta1 = np.reshape(nn_params[0:hidden_layer_size * (input_layer_size + 1)],
                        (hidden_layer_size, (input_layer_size + 1)),
↪order='F')
    Theta2 = np.reshape(nn_params[((hidden_layer_size * (input_layer_size +
↪1)))):],
                        (num_labels, (hidden_layer_size + 1)),
↪order='F')

    #Predicting
    pred = predict(Theta1, Theta2, X_test)
    pred = np.expand_dims(pred,axis=1)
    accuracy = (pred == Y_test).mean()*100
    accuracies[MaxIter] = accuracy
    pred = predict(Theta1, Theta2, X_train)
    pred = np.expand_dims(pred,axis=1)
    training_accuracy[MaxIter] = (pred == Y_train).mean()*100
    print('Testing Set Accuracy: ', accuracy)
    print('Training Set Accuracy: ', training_accuracy[MaxIter])
    print("Final Cost: ", cost[-1])
    cost_list.append(cost)

```

```
[67]: import matplotlib.pyplot as plt
xs, ys = zip(*accuracies.items())
plt.plot(xs, ys, "o-")
plt.xlabel("MaxIter")
plt.ylabel("Accuracy")
xs, ys = zip(*training_accuracy.items())
plt.plot(xs, ys, "ro-")
plt.legend(["testing_accuracy", "training_accuracy"])
```


[67]: <matplotlib.legend.Legend at 0x126b10910>



```
[68]: for cost_arr in cost_list:
    plt.plot(range(len(cost_arr)), cost_arr, "-")

plt.xlabel("iteration")
plt.ylabel("Cost")
plt.legend(["MaxIter = 100", "MaxIter = 150", "MaxIter = 200"])
```

[68]: <matplotlib.legend.Legend at 0x126b10100>



As one can clearly see in the above graphs, the NN achieves a higher accuracy for bigger MaxIter values, because it has more time to train. However, both accuracies increase less and less as MaxIter gets larger. So the training and testing accuracy's slope is anti-proportional to MaxIter, but the accuracies themselves are increasing together with MaxIter until they reach a point where the NN cannot fit the data better without overfitting (note that the lambda I chose for this experiment is  $\lambda = 7$ , so we do not overfit that easily). While the accuracy increases with the maximum number of iterations, the cost convergence is not as affected by it. You can clearly see the cost reaching a plateau early on, but the time when it reaches it is roughly the same for every value of MaxIter.

So, concludingly, one can say that lambda tries to prevent overfitting whilst sacrificing accuracy and cost-convergence and MaxIter increases the accuracies whilst also increasing the risk of overfitting the data. This means that a good model with nicely tuned hyperparameters should have reasonable values for lambda and MaxIter and, more importantly, a good tradeoff between their influences.

### 12.1.3 Improving on initial result from debugweights.mat

To improve, I simply loaded the debugweights into the cost function as initial parameters and just started optimizing from there. This model is probably highly overfit, since it has now an accuracy of 99.18% and the lambda value that was used is  $\lambda = 1$ . It was trained for 100 additional iterations.

```
[79]: print('Loading and Visualizing Data ...')

mat = scipy.io.loadmat('digitdata.mat')
X = mat['X']
```

```

y = mat['y']

print('Loading Saved Neural Network Parameters ...')

# Load the weights into variables Theta1 and Theta2
mat = scipy.io.loadmat('debugweights.mat');

# Unroll parameters
Theta1 = mat['Theta1']
Theta1_1d = np.reshape(Theta1, Theta1.size, order='F')
Theta2 = mat['Theta2']
Theta2_1d = np.reshape(Theta2, Theta2.size, order='F')

initial_nn_params = np.hstack((Theta1_1d, Theta2_1d))
MaxIter = 100
lambda_value = 1
# Create "short hand" for the cost function to be minimized
costFunction = lambda p : nnCostFunction(p, input_layer_size, hidden_layer_size,
                                         num_labels, X, y, lambda_value)

# Now, costFunction is a function that takes in only one argument (the
# neural network parameters)
[nn_params, cost] = fmincg(costFunction, initial_nn_params, MaxIter)
# Obtain Theta1 and Theta2 back from nn_params

Theta1 = np.reshape(nn_params[0:hidden_layer_size * (input_layer_size + 1)],
                    (hidden_layer_size, (input_layer_size + 1)),
                    order='F')
Theta2 = np.reshape(nn_params[(hidden_layer_size * (input_layer_size + 1)):],
                    (num_labels, (hidden_layer_size + 1)), order='F')

```

Loading and Visualizing Data ...

Loading Saved Neural Network Parameters ...

```

Iteration 1 | Cost: [0.3810212]
Iteration 2 | Cost: [0.37127893]
Iteration 3 | Cost: [0.36463237]
Iteration 4 | Cost: [0.36100255]
Iteration 5 | Cost: [0.35021463]
Iteration 6 | Cost: [0.34503721]
Iteration 7 | Cost: [0.34341822]
Iteration 8 | Cost: [0.34133862]
Iteration 9 | Cost: [0.33953897]
Iteration 10 | Cost: [0.33740099]
Iteration 11 | Cost: [0.33596374]
Iteration 12 | Cost: [0.33504536]
Iteration 13 | Cost: [0.33320868]
Iteration 14 | Cost: [0.33260408]
Iteration 15 | Cost: [0.33188161]

```

Iteration	16	Cost:	[0.33108026]
Iteration	17	Cost:	[0.33055596]
Iteration	18	Cost:	[0.3300339]
Iteration	19	Cost:	[0.3286211]
Iteration	20	Cost:	[0.32799557]
Iteration	21	Cost:	[0.32746564]
Iteration	22	Cost:	[0.32698966]
Iteration	23	Cost:	[0.32615938]
Iteration	24	Cost:	[0.32560775]
Iteration	25	Cost:	[0.32530679]
Iteration	26	Cost:	[0.32514242]
Iteration	27	Cost:	[0.3251085]
Iteration	28	Cost:	[0.3248847]
Iteration	29	Cost:	[0.32480782]
Iteration	30	Cost:	[0.32462838]
Iteration	31	Cost:	[0.32421904]
Iteration	32	Cost:	[0.32389556]
Iteration	33	Cost:	[0.32348843]
Iteration	34	Cost:	[0.32327297]
Iteration	35	Cost:	[0.32297332]
Iteration	36	Cost:	[0.32268525]
Iteration	37	Cost:	[0.32245862]
Iteration	38	Cost:	[0.32218229]
Iteration	39	Cost:	[0.32207838]
Iteration	40	Cost:	[0.32200989]
Iteration	41	Cost:	[0.32192025]
Iteration	42	Cost:	[0.32170581]
Iteration	43	Cost:	[0.32161984]
Iteration	44	Cost:	[0.32152271]
Iteration	45	Cost:	[0.32139461]
Iteration	46	Cost:	[0.32134078]
Iteration	47	Cost:	[0.32132509]
Iteration	48	Cost:	[0.32125513]
Iteration	49	Cost:	[0.3212128]
Iteration	50	Cost:	[0.32115362]
Iteration	51	Cost:	[0.32104472]
Iteration	52	Cost:	[0.32058975]
Iteration	53	Cost:	[0.32011243]
Iteration	54	Cost:	[0.32004432]
Iteration	55	Cost:	[0.32000438]
Iteration	56	Cost:	[0.31995882]
Iteration	57	Cost:	[0.31990999]
Iteration	58	Cost:	[0.31985374]
Iteration	59	Cost:	[0.31980389]
Iteration	60	Cost:	[0.31968306]
Iteration	61	Cost:	[0.31950328]
Iteration	62	Cost:	[0.31935435]
Iteration	63	Cost:	[0.31928132]

```
Iteration 64 | Cost: [0.31923547]
Iteration 65 | Cost: [0.31913285]
Iteration 66 | Cost: [0.31905641]
Iteration 67 | Cost: [0.31897811]
Iteration 68 | Cost: [0.31884922]
Iteration 69 | Cost: [0.31874822]
Iteration 70 | Cost: [0.31867737]
Iteration 71 | Cost: [0.31861595]
Iteration 72 | Cost: [0.31853479]
Iteration 73 | Cost: [0.3184923]
Iteration 74 | Cost: [0.31846892]
Iteration 75 | Cost: [0.31843724]
Iteration 76 | Cost: [0.31840971]
Iteration 77 | Cost: [0.31839738]
Iteration 78 | Cost: [0.31827248]
Iteration 79 | Cost: [0.31814848]
Iteration 80 | Cost: [0.31801761]
Iteration 81 | Cost: [0.3179069]
Iteration 82 | Cost: [0.31777663]
Iteration 83 | Cost: [0.31762719]
Iteration 84 | Cost: [0.31746506]
Iteration 85 | Cost: [0.31738721]
Iteration 86 | Cost: [0.31735524]
Iteration 87 | Cost: [0.31722898]
Iteration 88 | Cost: [0.31719164]
Iteration 89 | Cost: [0.31716785]
Iteration 90 | Cost: [0.31716145]
Iteration 91 | Cost: [0.31714139]
Iteration 92 | Cost: [0.31711694]
Iteration 93 | Cost: [0.31710172]
Iteration 94 | Cost: [0.31707391]
Iteration 95 | Cost: [0.31702016]
Iteration 96 | Cost: [0.31699543]
Iteration 97 | Cost: [0.31694357]
Iteration 98 | Cost: [0.3168766]
Iteration 99 | Cost: [0.31681225]
Iteration 100 | Cost: [0.31676604]
```

```
[80]: pred = predict(Theta1, Theta2, X)
      pred = np.expand_dims(pred,axis=1)
      print((pred == y).mean()*100)
```

99.18

**12.1.4 Imagine that you want to use a similar solution to classify 50x50 pixel grayscale images containing letters (consider an alphabet with 26 letters). Which changes would you need in the current code in order to implement this classification task?**

First of all, the input layer must be of size  $50 * 50 = 2500$  neurons. Then, the output layer should also resemble 26 classes instead of 10. This means that we would have way more parameters and a more complex problem, since our solution space became a lot bigger. So instead of  $401 * 26 + 26 * 10 = 10686$  parameters, we would now have to train  $2500 * 26 + 26 * 26 = 65676$  parameters, and that is only if we do not make any internal changes. Since the problem is more complex (we need to discriminate between more classes), we probably would also need additional internal neurons, maybe in different layers, which makes the optimization even more harder to solve. However, with convolutions, this could be achieved, since then we can efficiently decreased the amount of parameters to be trained because we would not have fully connected layers everywhere. Concretely, you would have to change `input_layer_size`, `hidden_layer_size` and `num_labels`, because based on these, the NN is built. Adding more layers or even convolutions with the current code would pose more of a challenge, because this version of the code is rather “hardcoded” to only have one hidden layer (we do not instantiate a list of weight matrices, but always use `Theta1` and `Theta2`, which makes for a big effort to change in every place we use it).

**12.1.5 How does your sigmoidGradient function work? Which is the return value for different values of z? How does it work with the input is a vector and with it is a matrix?**

Yes, it does, because numpy is doing the heavy load for us in this. See comments in the code at `sigmoidGradient.py` and `sigmoid.py`. I tested it with a couple of values and they all seem to fine.

**12.1.6 Change the value of the variable `show_examples` (in the python version, run the relevant block in the Jupyter one) in `ex_nn`, which information is provided? Did you get the expected information? Is anything unexpected there?**

Cannot find this variable in the code. I suspect the block where 100 examples are displayed is meant. When I change this value, less examples are shown, however the value must be a square, so 2, 4, 8, 16,... for the code to run properly. This is probably because only then it can be shown in a square format.